

- Name: Deep Salunkhe
- Roll No.:21102A0014
- [SEM-7 ML Lab5 Github Link](#)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

headers = ['ID', 'diagnosis', 'radius', 'texture', 'perimeter',
'area', 'smoothness', 'compactness', 'concavity', 'concave_points',
'symmetry', 'fractal_dimension',
          'radius2', 'texture2', 'perimeter2', 'area2',
'smoothness2', 'compactness2', 'concavity2', 'concave_points2',
'symmetry2', 'fractal_dimension2',
          'radius3', 'texture3', 'perimeter3', 'area3',
'smoothness3', 'compactness3', 'concavity3', 'concave_points3',
'symmetry3', 'fractal_dimension3']
data = pd.read_csv("/content/sample_data/wdbc.data", delimiter="," ,
names = headers)
```

data

```
{"type": "dataframe", "variable_name": "data"}
```

data.shape

```
(569, 32)
```

data.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	569 non-null	int64
1	diagnosis	569 non-null	object
2	radius	569 non-null	float64
3	texture	569 non-null	float64
4	perimeter	569 non-null	float64
5	area	569 non-null	float64
6	smoothness	569 non-null	float64
7	compactness	569 non-null	float64
8	concavity	569 non-null	float64
9	concave_points	569 non-null	float64
10	symmetry	569 non-null	float64
11	fractal_dimension	569 non-null	float64
12	radius2	569 non-null	float64

13	texture2	569	non-null	float64
14	perimeter2	569	non-null	float64
15	area2	569	non-null	float64
16	smoothness2	569	non-null	float64
17	compactness2	569	non-null	float64
18	concavity2	569	non-null	float64
19	concave_points2	569	non-null	float64
20	symmetry2	569	non-null	float64
21	fractal_dimension2	569	non-null	float64
22	radius3	569	non-null	float64
23	texture3	569	non-null	float64
24	perimeter3	569	non-null	float64
25	area3	569	non-null	float64
26	smoothness3	569	non-null	float64
27	compactness3	569	non-null	float64
28	concavity3	569	non-null	float64
29	concave_points3	569	non-null	float64
30	symmetry3	569	non-null	float64
31	fractal_dimension3	569	non-null	float64

dtypes: float64(30), int64(1), object(1)

memory usage: 142.4+ KB

```
pip install xgboost
```

Requirement already satisfied: xgboost in

/usr/local/lib/python3.10/dist-packages (2.1.1)

Requirement already satisfied: numpy in

/usr/local/lib/python3.10/dist-packages (from xgboost) (1.26.4)

Requirement already satisfied: nvidia-nccl-cu12 in

/usr/local/lib/python3.10/dist-packages (from xgboost) (2.23.4)

Requirement already satisfied: scipy in

/usr/local/lib/python3.10/dist-packages (from xgboost) (1.13.1)

```
from sklearn.model_selection import train_test_split
```

```
X = data.drop(['diagnosis', 'ID'], axis=1)
```

```
# Map 'B' to 0 and 'M' to 1
```

```
y = data['diagnosis'].map({'B': 0, 'M': 1})
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

RANDOM FOREST

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_classifier = RandomForestClassifier()
```

```

rf_classifier.fit(X_train, y_train)

y_pred = rf_classifier.predict(X_test)

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, precision_score, recall_score, f1_score, roc_curve,
auc

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, pos_label=1)
recall = recall_score(y_test, y_pred, pos_label=1)
f1 = f1_score(y_test, y_pred, pos_label=1)

print(f'Accuracy: {accuracy}')
print(f'Precision: {precision}')
print(f'Recall: {recall}')
print(f'F1-score: {f1}')

print(classification_report(y_test, y_pred))

print(confusion_matrix(y_test, y_pred))

Accuracy: 0.9649122807017544
Precision: 0.975609756097561
Recall: 0.9302325581395349
F1-score: 0.9523809523809523

```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	71
1	0.98	0.93	0.95	43
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

```

[[70  1]
 [ 3 40]]

y_probs = rf_classifier.predict_proba(X_test)[:, 1]

fpr, tpr, thresholds = roc_curve(y_test, y_probs, pos_label=1)


auc_value = auc(fpr, tpr)
print(f'AUC: {auc_value:.4f}')

plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, marker='.', label=f'ROC curve (AUC =
{auc_value:.4f})')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')

```

```
plt.legend()
plt.grid()
plt.show()
```

AUC: 0.9920



XGBOOST

```
from xgboost import XGBClassifier

# Initialize XGBoost Classifier
xgb_classifier = XGBClassifier()

# Fit the model
xgb_classifier.fit(X_train, y_train)

# Make predictions
y_pred = xgb_classifier.predict(X_test)

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, precision_score, recall_score, f1_score, roc_curve,
auc

accuracy = accuracy_score(y_test, y_pred)
```

```
precision = precision_score(y_test, y_pred, pos_label=1)
recall = recall_score(y_test, y_pred, pos_label=1)
f1 = f1_score(y_test, y_pred, pos_label=1)
```

```
print(f'Accuracy: {accuracy}')
print(f'Precision: {precision}')
print(f'Recall: {recall}')
print(f'F1-score: {f1}')
```

```
print(classification_report(y_test, y_pred))
```

```
print(confusion_matrix(y_test, y_pred))
```

```
Accuracy: 0.956140350877193
Precision: 0.9523809523809523
Recall: 0.9302325581395349
F1-score: 0.9411764705882353
```

	precision	recall	f1-score	support
0	0.96	0.97	0.97	71
1	0.95	0.93	0.94	43

accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

```
[[69  2]
 [ 3 40]]
```


```
y_probs = xgb_classifier.predict_proba(X_test)[:, 1]
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_probs, pos_label=1)
```

```
auc_value = auc(fpr, tpr)
print(f'AUC: {auc_value:.4f}')
```

```
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, marker='.', label=f'ROC curve (AUC = {auc_value:.4f})')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.grid()
plt.show()
```

```
AUC: 0.9908
```



#Regression

```
data = pd.read_csv("/content/sample_data/housing.csv")
data

{"summary": "{\n  \"name\": \"data\",\n  \"rows\": 20640,\n  \"fields\": [\n    {\n      \"column\": \"longitude\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2.0035317235025882,\n        \"min\": -124.35,\n        \"max\": -114.31,\n        \"num_unique_values\": 844,\n        \"samples\": [\n          -118.63,\n          -119.86,\n          -121.26\n        ],\n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"latitude\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2.1359523974571153,\n        \"min\": 32.54,\n        \"max\": 41.95,\n        \"num_unique_values\": 862,\n        \"samples\": [\n          33.7,\n          34.41,\n          38.24\n        ],\n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"housing_median_age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 12.58555761211165,\n        \"min\": 1.0,\n        \"max\": 52.0,\n        \"num_unique_values\": 52,\n        \"samples\": [\n          25.0,\n          7.0\n        ],\n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"total_rooms\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 12.58555761211165,\n        \"min\": 1.0,\n        \"max\": 52.0,\n        \"num_unique_values\": 52,\n        \"samples\": [\n          25.0,\n          7.0\n        ],\n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }\n  ]\n}
```


#	Column	Non-Null	Count	Dtype
0	longitude	20640	non-null	float64
1	latitude	20640	non-null	float64
2	housing_median_age	20640	non-null	float64
3	total_rooms	20640	non-null	float64
4	total_bedrooms	20433	non-null	float64
5	population	20640	non-null	float64
6	households	20640	non-null	float64
7	median_income	20640	non-null	float64
8	median_house_value	20640	non-null	float64
9	ocean_proximity	20640	non-null	object

dtypes: float64(9), object(1)

memory usage: 1.6+ MB

data.isnull().sum()

longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	207
population	0
households	0
median_income	0
median_house_value	0
ocean_proximity	0

dtype: int64

data.dropna(inplace=True)

data.isnull().sum()

longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	0
population	0
households	0
median_income	0
median_house_value	0
ocean_proximity	0

dtype: int64

data.reset_index(inplace=True,drop=True)

data['ocean_proximity'].value_counts()

ocean_proximity	
<1H OCEAN	9034


```
INLAND      6496
NEAR OCEAN  2628
NEAR BAY    2270
ISLAND      5
Name: count, dtype: int64
```

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
data['ocean_proximity']=le.fit_transform(data['ocean_proximity'])
```

```
data["rooms_per_household"] = data["total_rooms"]/data["households"]
data["bedrooms_per_room"] = data["total_bedrooms"]/data["total_rooms"]
data["population_per_household"]=data["population"]/data["households"]
```

```
data.corr()
```

```
{"summary":{"\n  \"name\": \"data\", \n  \"rows\": 13, \n  \"fields\": [\n    {\n      \"column\": \"longitude\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.4064425796486295, \n        \"min\": -0.9246161131160101, \n        \"max\": 1.0, \n        \"num_unique_values\": 13, \n        \"samples\": [\n          0.09265683306977554, \n          -0.28953010139097807, \n          1.0\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }, \n    {\n      \"column\": \"latitude\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.4047506201096927, \n        \"min\": -0.9246161131160101, \n        \"max\": 1.0, \n        \"num_unique_values\": 13, \n        \"samples\": [\n          -0.11381506848357399, \n          0.2008010727260355, \n          -0.9246161131160101\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }, \n    {\n      \"column\": \"housing_median_age\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.3533184865983313, \n        \"min\": -0.36062829984244227, \n        \"max\": 1.0, \n        \"num_unique_values\": 13, \n        \"samples\": [\n          0.13608923857356492, \n          0.11233014464562725, \n          -0.10935654863027307\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }, \n    {\n      \"column\": \"total_rooms\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.474090511200691, \n        \"min\": -0.36062829984244227, \n        \"max\": 1.0, \n        \"num_unique_values\": 13, \n        \"samples\": [\n          -0.18790000413461336, \n          0.015363166414694703, \n          0.0454801674218395\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }, \n    {\n      \"column\": \"total_bedrooms\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.47841935648102935, \n        \"min\": -0.32045104175060396, \n        \"max\": 1.0, \n        \"num_unique_values\": 13, \n        \"samples\": [\n          0.08423813762384522, \n          0.014767943833213214, \n          0.06960802175408133\n        ], \n
```

```

\"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"population\", \n    \"properties\": { \n      \"dtype\": \"number\", \n      \"std\": 0.4678296793254769, \n      \"min\": -0.2957872971044803, \n      \"max\": 1.0, \n      \"num_unique_values\": 13, \n      \"samples\": [ \n        0.03531932605444879, \n        0.06962966726072072, \n        0.10027030094083503 \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"households\", \n    \"properties\": { \n      \"dtype\": \"number\", \n      \"std\": 0.48289530521105845, \n      \"min\": -0.30276797344891637, \n      \"max\": 1.0, \n      \"num_unique_values\": 13, \n      \"samples\": [ \n        0.06508685650819257, \n        0.018250608808529776, \n        0.056512772430637834 \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"median_income\", \n    \"properties\": { \n      \"dtype\": \"number\", \n      \"std\": 0.3946564083548571, \n      \"min\": -0.6156606088731494, \n      \"max\": 1.0, \n      \"num_unique_values\": 13, \n      \"samples\": [ \n        -0.6156606088731494, \n        0.014678593813623618, \n        -0.015550150379729375 \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"median_house_value\", \n    \"properties\": { \n      \"dtype\": \"number\", \n      \"std\": 0.33973304362366785, \n      \"min\": -0.25588014941949866, \n      \"max\": 1.0, \n      \"num_unique_values\": 13, \n      \"samples\": [ \n        -0.25588014941949866, \n        0.08048786002048676, \n        -0.04539821933443104 \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"ocean_proximity\", \n    \"properties\": { \n      \"dtype\": \"number\", \n      \"std\": 0.2997284541302523, \n      \"min\": -0.28953010139097807, \n      \"max\": 1.0, \n      \"num_unique_values\": 13, \n      \"samples\": [ \n        0.002106848712775119, \n        -0.28953010139097807 \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"rooms_per_household\", \n    \"properties\": { \n      \"dtype\": \"number\", \n      \"std\": 0.32868946618396233, \n      \"min\": -0.41695223734049447, \n      \"max\": 1.0, \n      \"num_unique_values\": 13, \n      \"samples\": [ \n        -0.0026413047510256043, \n        0.41695223734049447, \n        0.027306806809850443 \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  }, \n  { \n    \"column\": \"bedrooms_per_room\", \n    \"properties\": { \n      \"dtype\": \"number\", \n      \"std\": 0.3768344965853683, \n      \"min\": -0.6156606088731494, \n      \"max\": 1.0, \n      \"num_unique_values\": 13, \n      \"samples\": [ \n        1.0, \n        0.002106848712775119, \n        0.09265683306977554 \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  } \n}

```

```

n    },\n    {\n        \"column\": \"population_per_household\", \n        \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.27877451835446043, \n            \"min\": -0.02835486367066539, \n            \"max\": 1.0, \n            \"num_unique_values\": 13, \n            \"samples\": [\n                0.00293785249526524, \n                0.009151877026194142, \n                0.002303875515041799\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\"\n        }\n    }\n    ],\n    \"type\": \"dataframe\"}

```

```

x_data = data.drop([\"median_house_value\"], axis = 1).values
y_data = data[\"median_house_value\"].values

```

```

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_data, y_data,
                                                    test_size=0.25, random_state=42)

```

RANDOM FOREST

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

```

```
rf_regressor = RandomForestRegressor()
```

```
rf_regressor.fit(x_train, y_train)
```

```
y_pred = rf_regressor.predict(x_test)
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f'Mean Squared Error: {mse}')
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f'R-squared: {r2}')
```

```
Mean Squared Error: 2631339567.8051205
```

```
R-squared: 0.8055012570233084
```

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(10, 6))
```

```
plt.scatter(y_test, y_pred, alpha=0.5, edgecolor='k')
```

```
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
         'r--', lw=2)
```


```
plt.xlabel('Actual Values')
```

```
plt.ylabel('Predicted Values')
```

```
plt.title('Actual vs Predicted Values')
```

```
plt.grid(True)
```

```
plt.show()
```



XGBOOST

```
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Initialize XGBoost Regressor
xgb_regressor = XGBRegressor()

# Fit the model
xgb_regressor.fit(x_train, y_train)

# Make predictions
y_pred = xgb_regressor.predict(x_test)

from sklearn.metrics import mean_squared_error, r2_score

mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse}')

r2 = r2_score(y_test, y_pred)
print(f'R-squared: {r2}')

Mean Squared Error: 2161493666.7455506
R-squared: 0.8402305022590588
```

```
# Randomly sample 1000 points or fewer if your data has fewer points
sample_size = min(1000, len(y_test))
indices = np.random.choice(len(y_test), size=sample_size,
replace=False)

# Sampled data
y_test_sample = y_test[indices]
y_pred_sample = y_pred[indices]

# Create the plot
plt.figure(figsize=(10, 6))

# Scatter plot
plt.scatter(y_test_sample, y_pred_sample, alpha=0.6, edgecolor='k',
s=50, color='skyblue', label='Predictions')

# Line of perfect prediction
plt.plot([y_test_sample.min(), y_test_sample.max()],
[y_test_sample.min(), y_test_sample.max()], 'r--', lw=2,
label='Perfect Prediction')

# Adding labels and title
plt.xlabel('Actual Values', fontsize=14)
plt.ylabel('Predicted Values', fontsize=14)
plt.title('Actual vs Predicted Values (Sampled)', fontsize=16)
plt.legend()
plt.grid(True, linestyle='--', alpha=0.7)

# Show the plot
plt.tight_layout()
plt.show()
```

